

Synchronization

Concurrent Processes

- # Concurrent processes can:
 - ❖ Compete for shared resources
 - ❖ Cooperate with each other in sharing the global resources
- # OS deals with competing processes
 - ❖ carefully allocating resources
 - ❖ properly isolating processes from each other
- # OS deals with cooperating processes
 - ❖ providing mechanisms to share resources

Processes: *Competing*

- Processes that do not work together cannot affect the execution of each other, but they can compete for devices and other resources
- **Example:** Independent processes running on a computer
- **Properties:** Deterministic; reproducible
 - ◆ Can stop and restart without side effects
 - ◆ Can proceed at arbitrary rate

Processes: Cooperating

- Processes that are aware of each other, and directly (by exchanging messages) or indirectly (by sharing a common object) work together, may affect the execution of each other
- **Example:** Transaction processes in an airline reservation system
- **Properties:**
 - ◆ Share a common object or exchange messages
 - ◆ Non-deterministic; May be irreproducible
 - ◆ Subject to *race conditions* – coming up!

Why Cooperation?

- Cooperation clearly presents challenges – why do we want it?
- We may want to share resources:
 - ◆ Sharing a global task queue
- We may want to do things faster:
 - ◆ Read next block while processing current one; divide a job into pieces and execute concurrently
- We may want to solve problems modularly

UNIX example:

```
cat infile | tr ' ' '\012' | tr '[A-Z]' '[a-z]' | sort | uniq -c
```

Problem with

Original application..

Correct output is always the serial one

- Hard to understand or executing programs
- Instructions of the program are arbitrary

$$\begin{array}{l} A = 1 \\ B = 2 \\ A = B + 1 \\ B = B * 2 \end{array}$$

order is important

shared memory, a, b same address between two instances

- ◆ For cooperating processes, the order of (some) instructions are irrelevant
- ◆ However, certain instruction combinations must be avoided, for example:

Process A	Process B	concurrent access
$A = 1;$	$B = 2;$	<i>does not matter</i>
$A = B + 1;$	$B = B * 2;$	<i>important!</i>

Race Conditions

- When two or more processes are reading or writing some shared data and final result depends on who runs precisely when is a *race condition*
- How to avoid race conditions?
 - ◆ Prohibit more than one process from reading and writing shared data at the same time
 - ◆ Essentially we need *mutual exclusion* 
- Mutual exclusion:
 - ◆  When one process is reading or writing a shared data, other processes should be prevented from doing the same
 - ◆ Providing mutual exclusion in OS is a major design issue

Suppose two processes A & B are updating a shared account. A withdrawing and B depositing. Initial balance = \$100

An Example

Correct balance = \$100
(\$50 withdrawal and \$50 deposit)

Initial balance = \$100

Process A

```
R1 <- $100
R2 <- $50
R1 <- $100 - $50
```

balance <- R1 = \$50
store R1, balance

balance <- R3 = \$150
store R3, balance

Process A
load R1, balance
load R2, 50
sub R1, R2
store R1, balance

Process B
load R3, balance
load R4, 50
add R3, R4
store R3, balance

Possible values for balance are \$50, \$100, and \$150!

Process B

```
R3 <- $100
R4 <- $50
R3 <- $100 + $50
```

Possible values for balance are \$50, \$150, and \$100!

An Example....

- Suppose we have the following code for the account transactions

```
authenticate user  
open account  
load R1, balance  
load R2, amount (-ve for withdrawal)  
add R1, R2  
store R1, balance  
close account  
display account info
```

An Example....

- Suppose we have the following code for the account transactions

authenticate user

open account

load R1, balance

load R2, amount (-ve for withdrawal)

add R1, R2

store R1, balance

close account

display account info

another program have to wait to access the critical section

Critical section

lock

Critical Section

- Part of the program that accesses shared data
- If we arrange the program runs such that no two ~~programs~~ are in their critical sections at the ~~same~~ time, race conditions can be avoided
 - processes

Critical Section

For multiple programs to cooperate correctly and efficiently:

- No two processes may be **simultaneously** in their critical sections
- No assumptions be made about **speeds** or number of CPUs
- No process running **outside** its critical section may block other processes
- No process should have to wait **forever** to enter its critical section

Requirements of a Critical Section

Critical Sections

- Critical region execution should be *atomic* i.e., its runs “all or none”
- Should not be interrupted in the middle
- For efficiency, we need to minimize the length of the critical sections
- Most inefficient scenario
 - ◆ Whole program is a critical section – no multiprogramming!

Road to a Solution

Simplest solution:

- ❖ Disable all interrupts
 - ❖ This should work in a single processor situation
 - ❖ Because interrupts cause out-of-order execution due to interrupt servicing
 - ❖ With interrupts disabled, a process will run until it yields the resource voluntarily
- ❖ Not practical:
 - ❖ OS won't allow a user process to disable all interrupts – OS operation will be hindered too!
 - ❖ Does not work on multiprocessors

only sleeping beauty approach, no alarm approach

Road to a Solution

- Idea: To come up with a **software solution**
 - hardware doesn't allow to disable interrupts
- Use “lock” variables to prevent two processes entering the critical section at the same time – all along we are talking about a single critical section
 - Use variable lock
 - If `lock == 0` set `lock = 1` and `enter_region`
 - If `lock == 1` wait until `lock becomes 0`
 - Does not work, Why??

They have to alternate. If one process doesn't go into the critical again -> pb

Strict Alternation

alternation solve previous pb

violate condition 3

- Two processes take turns in entering the critical section

- Global variable turn set either to 0 or 1

Process 0

Process 1

```
while (TRUE) {           while (TRUE) {  
    while (turn !=0);      while (turn !=1);  
    critical_section();    critical_section();  
    turn = 1;              turn = 0;  
    non_critical();        non_critical();  
}
```

Strict Alternation

- What is the problem with strict alternation?
- We can have starvation, why??
- What are the other drawbacks?
 - ◆ Continuously testing a variable until **some** value appears is called **busy waiting**
 - ◆ Busy waiting wastes CPU time – should be used when the expected wait is short
 - ◆ A lock that uses busy waiting is called **spin lock**

one process never have the chance to
get in the critical section

Example

Quantum : The Duration of time slice

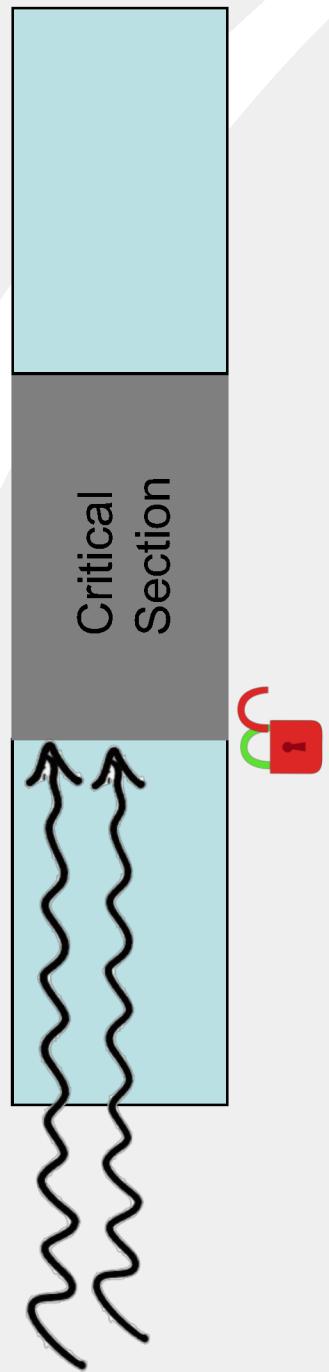
Threads are used for concurrent processing in a data-parallel application. The application has a critical section. If a single thread of execution is used the critical section takes 2 seconds and other parts take 6 seconds. We have a large dataset so we use 2 threads. What is the run time? We have an even larger dataset so we use 8 threads. What is the run time?

one thread $6 + 2 = 8$ s
two thread 4 sec for cs + $2 + 6 + 6 = 18$
W busy waiting, we lost 2 sec.
Use sleep instead
 $16.5 (.5 \text{ is the sleep wake up overhead})$

Key

Assume single CPU and ideal time sharing (no overhead). All threads starting at the same time.

Locks – An Illustration



Mutual Exclusion: First Attempt

- The simplest mutual exclusion strategy is taking turns as considered earlier

```
/* process 0 */  
.  
.  
while (turn != 0);  
/* critical section */  
turn = 1;  
.
```

turn strategy - not working solution

```
/* process 1 */  
.  
.  
while (turn != 1);  
/* critical section */  
turn = 0;  
.
```

Mutual Exclusion: Second Attempt

- First attempt problem – single key shared by the two processes
- Each have their own key to the critical section
- Solution does not work! Why??

```
/* process 0 */  
.  
.  
.  
if process want to go in CS set flag to true  
while (flag[1]);  
flag[0] = true;  
/* critical section */  
flag[0] = false;  
  
Pb : no lock for the lock  
.
```

```
/* process 1 */  
.  
.  
while (flag[0]);  
flag[1] = true;  
/* critical section */  
flag[1] = false;  
.
```

Mutual Exclusion: Third Attempt

- This works, i.e., provides mutual exclusion
- Has **deadlock** – why?

[Both will be locked out of the Cs if context switch happens in the middle of the lock writing.]

```
/* process 1 */  
.  
. flag[1] = true;  
while (flag[0]);  
/* critical section */  
flag[1] = false;  
.  
.  
.
```



```
/* process 0 */  
.  
. flag[0] = true;  
while (flag[1]);  
/* critical section */  
flag[0] = false;  
.  
.  
.
```

Mutual Exclusion: Fourth Attempt

- This works
- Has livelock – why?

If delay not random

```
/* process 0 */  
.  
.  
flag[0] = true;  
while (flag[1])  
{  
    flag[0] = false;  
    /* random delay */  
    flag[0] = true;  
}  
/* critical section */  
flag[0] = false;  
.  
.
```

```
/* process 1 */  
.  
.  
flag[1] = true;  
while (flag[0])  
{  
    flag[1] = false;  
    /* random delay */  
    flag[1] = true;  
}  
/* critical section */  
flag[1] = false;  
.  
.
```

Peterson's Algorithm

work only for 2 processes

exists a generalization for n processes

use busy waiting

```
/* process 0 */  
...  
flag[0] = true;  
turn = 1; give priority,  
while (flag[1] &&  
        turn == 1);  
/* critical section */  
flag[0] = false;  
/* remainder */  
...
```

```
/* process 1 */  
...  
flag[1] = true;  
turn = 0;  
while (flag[0] &&  
        turn == 0);  
/* critical section */  
flag[1] = false;  
/* remainder */  
...
```

On the user level,

Can Hardware Help?

- Dekker's and Peterson's algorithms are pure software solutions
- Can hardware provide any help?
 - ◆ make the solution more efficient
 - ◆ make it scalable to more processes?
- Yes!
- Current microprocessors have hardware instructions supporting mutual exclusion

Test and Lock

works for any number of processes

- TSL RX, LOCK is a typical CPU instruction providing support for mutual exclusion
 - ◆ read the contents of memory location LOCK into register RX and stores a non-zero value at memory location LOCK
 - ◆ operation of reading and writing are **indivisible - atomic**

enter_section:	TSL REGISTER, LOCK memory location	move lock to reg	lock was zero?	get 0, can proceed in the CS
	// copy lock to reg and set lock to 1	CMP REGISTER, #0	JNE enter_section	At the same time read and write to 1
		// was lock zero??	// if non zero, lock was set	
		JNE enter_section	RET	

can be tricky on multiprocessors, 2 cores can trigger the same atomic instruction at the same time

Properties of Machine Instruction Approach

Advantages:

- ◆ applicable to any number of processes
- ◆ can be used with single processor or multiple processors that **share a single memory**
- ◆ simple and easy to verify
- ◆ can be used to support multiple critical sections, i.e., define a separate variable for each critical section

how we prevent
multiple TSL Rx
sent at the same
time

Properties of Machine Instruction Approach

- Disadvantages:
 - ◆ **Busy waiting is employed** – process waiting to get into a critical section consumes CPU time
 - ◆ **Starvation is possible** – selection of entering process is arbitrary when multiple processes are contending to enter

Busy Waiting Approaches

- Mutual exclusion schemes discussed so far are based on busy waiting
- Busy waiting not desirable:
- Suppose a computer runs two processes H: high priority and L: low priority
 - ◆ scheduler always runs H when it is in ready state
 - ◆ at a certain time L is in its critical section and H become runnable
 - ◆ H begins busy waiting to enter the critical section
 - ◆ L is never scheduled to leave the critical section
 - ◆ there is a **deadlock**
 - ◆ this situation is sometimes referred to as the **priority inversion problem**

Sleep/Wakeup Approach

- Alternative to busy waiting that is inefficient and deadlock prone is to use a sleep/wakeup approach
- Implemented by the Semaphores

Semaphores

Fundamental principle:

- ❖ two or more processes can cooperate by sending simple messages
- ❖ a process can be forced to stop at a specific place until it receives a specific message
- ❖ complex coordination can be satisfied by appropriately structuring these messages
- ❖ for messaging a special variable called **semaphore s** is used
- ❖ to transmit a message via a semaphore a process executes signal(s)
- ❖ to receive a message via a semaphore a process executes wait(s)

Semaphores

smaller number has higher priority

avoid busy waiting



- Operations defined on a semaphore:
 - can be initialized to a nonnegative value – set semaphore

- wait – decrements the semaphore value – if value becomes negative, process executing wait is blocked
 - status changes - go to sleep
lock is not available when value negative
- signal – increments the semaphore value – if value is not positive, a process blocked by a wait operation is unblocked

Sleep gives overhead

we can have a queue of waiting processes

file descriptor table isn't process private data

Semaphores

A definition of semaphore primitives...

```
struct semaphore {
    int count;                                how many processes are waiting
    queueType queue;                          contains waiting processes
}

void signal(semaphore s)
{
    s.count++;                                if (s.count <= 0)
    {                                         remove a process P
        if (s.count < 0)                      from s.queue;
        {                                     place process P on
            place this process in           ready list;
            s.queue;                         }
        block this process
    }
}
```

```
void wait(semaphore s)
{
    s.count--;                                if (s.count < 0)
    {                                         place this process in
        if (s.count < 0)                      s.queue;
        {                                     block this process
            place this process in           }
        s.queue;                         }
    }
}
```

data structure created inside the kernel

Semaphores

- Wait and Signal primitives are assumed to be atomic
 - ◆ they cannot be interrupted and treated as an indivisible step
- A queue is used to hold the processes waiting on a semaphore
 - How are the processes removed from this queue?
 - ◆ FIFO: process blocked longest should be released next – **strong semaphore**
 - ◆ Order not specified – **weak semaphore**

Type of semaphore

- Semaphores are usually ways:
 - **binary**—is a semaphore with of 0 or 1

These are initial values.
A semaphore can have negative value at any time .. many threads or processes waiting on it

- Or a value of FALSE (TRUE if you prefer
 - Initialized to 1 for mutex applications
- **counting**—is a semaphore with an integer value ranging between 0 and an arbitrarily large number - initial value might represent the number of units of the critical resources that are available - also known as a **general** semaphore

Producer-Consumer: Semaphores

avoid busy waiting

shared buffer

buffer can become full

```
semaphore mutex = 1;  
semaphore empty = N;  
semaphore full = 0;
```

counting semaphore

```
// protects the critical section  
// counts the empty slots  
// counts full buffer slots
```

N = max size of buffer

```
producer() {  
    int item;  
  
    while(TRUE) {  
        item = produce_item();  
        wait(&empty);  
        wait(&mutex);  
        insert_item(item);  
        signal(&mutex);  
        signal(&full);  
    }  
}
```

C

only wait if no empty

Mutex is a semaphore by

increment full

signal tell the consumer updates happened,

semaphore can only be initialized to positive values

```
consumer() {  
    int item;
```

```
    while(TRUE) {  
        wait(&full);  
        wait(&mutex);  
        item = remove_item();  
        signal(&mutex);  
        signal(&empty);  
        consume_item(item);  
    }  
}
```

wait until there is smt to read

if smt waiting, semaphore should be neg. Otherwise positive.

Producer-Consumer: Semaphores

underunning, overrunning

- One binary semaphore **mutex** to ensure only one process is manipulating the buffers at a time
- **Empty** and **Full** counting semaphores, to count the number of empty or full buffer slots respectively
 - ◆ Empty initialized to N, so waiting on empty means waiting if there are no empty slots (*buffer full!*)
 - ◆ Full initialized to 0, so waiting on full means waiting if there are no full slots (*buffer empty!*)

Primitives Revisited

- The synchronization primitives discussed so far are all of the form:

```
entry protocol
```

```
< access data >
```

```
exit protocol
```

- Semaphores give us some abstraction: implementation of the protocol to access shared data is transparent to the user
 - More abstraction would be of additional benefit (as it often is!)

Monitors

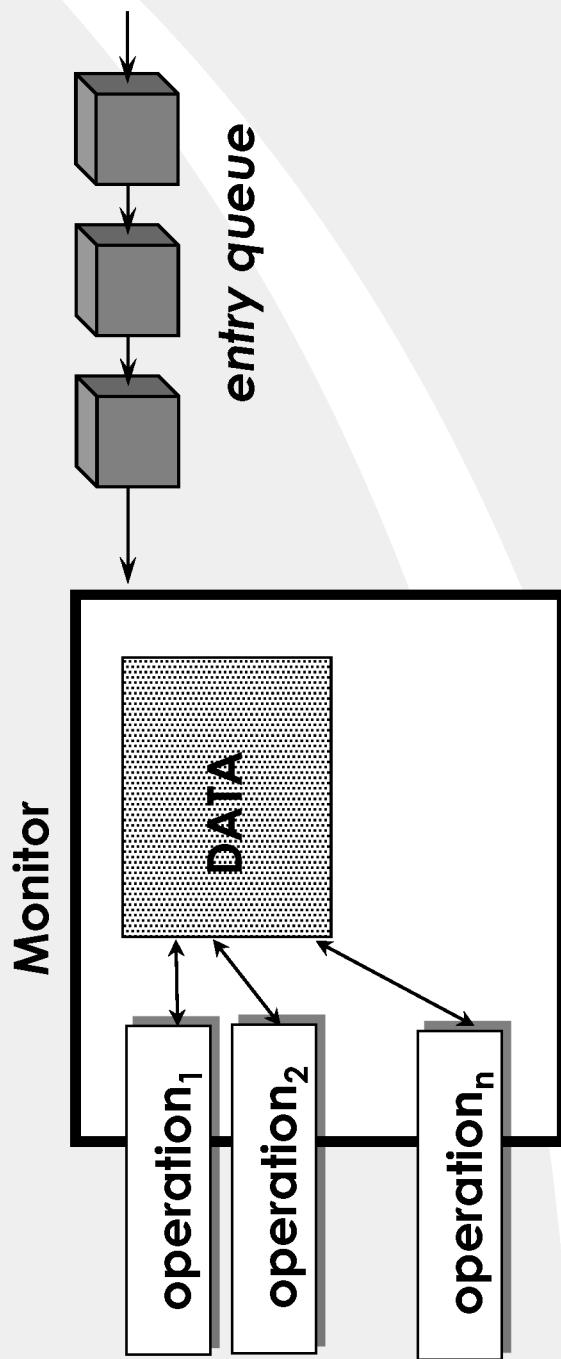
- A monitor is a high-level (programming language) abstraction that combines (and hides) the following:
 - shared data
 - operations on the data
 - synchronization using condition variables

Monitors

- Mutual exclusion --- not the only thing we need for concurrency
- Abstraction provided by monitors allows us to deal with many different issues in concurrency
 - ◆ block processes when a resource is busy
- With a monitor data type, we define a set of variables that define the state of an instance of the type, and a set of procedures that define the externally available operations on the type

Monitors

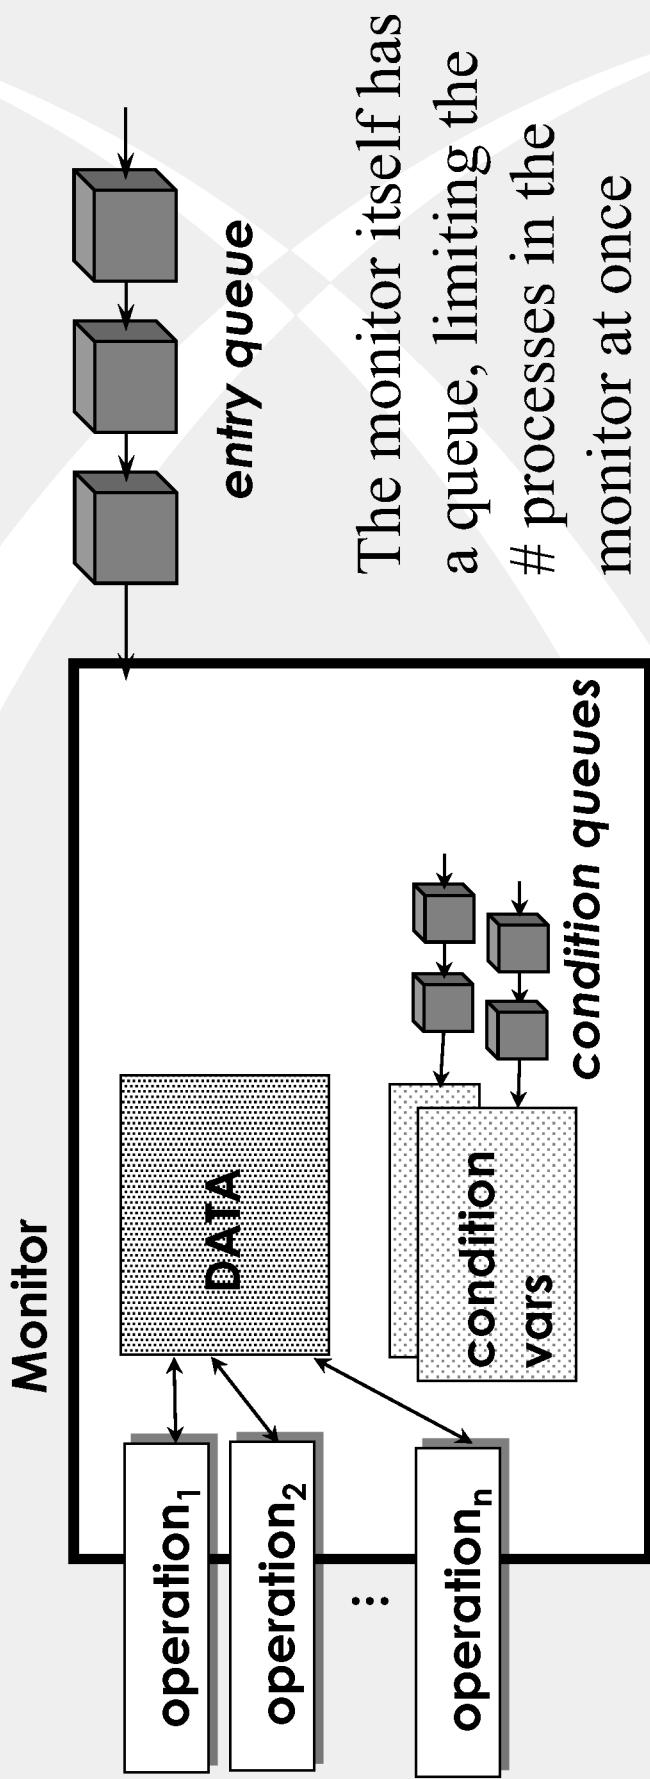
- Monitors are more than just objects
- A monitor ensures that only one process at a time can be active within the monitor – so you don't need to code this explicitly



Monitors

- Just this on its own though isn't enough – need the additional synchronization mechanisms
- Monitors use **condition variables** to provide user-tailored synchronization and manage each with a separate queue for each condition
- The only operations available on these variables are **WAIT** and **SIGNAL**
 - ◆ wait suspends process that executes it until someone else does a **signal** (different from semaphore wait!)
 - ◆ signal resumes one suspended process. no effect if nobody's waiting (different from semaphore free/signal)!

Monitor Abstraction



How Monitors Work

- Remember only one process can be active in the monitor at once. Some will be waiting to get in, others will be waiting on conditions in the Hoare's Monitor
- If P executes x.signal (signal for condition x), and there's a process Q suspended on x, both can't be active at once. 2 possibilities:
 - ◆ P waits on some condition till Q leaves monitor
 - ◆ Q waits on some condition till P leaves monitor
- **P executes signal as the last operation, it leaves monitor automatically, resuming Q**

Problems with *synch primitives*

- *Starvation*: the situation in which some processes are making progress toward completion but some others are locked out of the resource(s)
- *Deadlock*: the situation in which two or more processes are locked out of the resource(s) that are held by each other

Problems with *synch* Primitives

- The most important deficiency of the primitives discussed so far is that they were all designed on the concept of one or more CPUs accessing a ‘common’ memory
- Hence, these primitives are not applicable to distributed systems. **Solution?** Message passing...

Synch Primitives—Summary

