

# COMP 551 - Applied Machine Learning

## Lecture 2 - Linear Regression

---

William L. Hamilton

(with slides and content from Joelle Pineau)

\* Unless otherwise noted, all material posted for this course are copyright of the instructor(s), and cannot be reused or reposted without the instructor's written permission.

# Some logistics (revisited)

---

- Course website: <https://www.cs.mcgill.ca/~wlh/comp551/>
  - Used for: Schedule, **links to lectures slides**, homework materials, and readings
  - TA office hours are up on the website.
- MyCourses
  - Used for: Announcements, quizzes, grading, and discussions.
  - **There is now a discussion forum for finding project partners.**

# Quizzes

---

- There will be two quizzes a week (starting this week).
  - One available Tuesday-Wednesday.
  - The other available Thursday-Friday.
- They will be on MyCourses.
- The two weekly quizzes will be similar, but not identical.
- Only your best quiz from each week will count towards your grade.
- Important: This week's quizzes are for practice and self-assessment. They will not count towards your grade.

# Python tutorial hours

---

- There will be Python tutorial hours in Trottier 3120:
  - 6-8pm on Friday, Sept. 13th
  - 4:30-6:30pm on Monday, Sept. 16<sup>th</sup>.
  - 3-5 pm on Tuesday, Sept. 17<sup>th</sup>.
  - 4-6 pm on Wednesday, Sept. 18<sup>th</sup>.
- Come if you have never used Python and need help!
- There is limited space and they will be first-come-first-served.
- Please attend the whole session if you go.

# Supervised learning

- Given a set of training examples:  $\mathbf{x}_i = \langle x_{i,1}, x_{i,2}, x_{i,3}, \dots, x_{i,m}, y_i \rangle$   
 $x_{i,j}$  is the  $j^{\text{th}}$  feature of the  $i^{\text{th}}$  example  
 $y_i$  is the desired output (or target) for the  $i^{\text{th}}$  example.
- We want to learn a function  $f: X_1 \times X_2 \times \dots \times X_m \rightarrow Y$ , which maps the input variables to the output/target.
- Formally,  $f$  is called the hypothesis (or model).

tumor size	texture	perimeter	shade	outcome	size change
18.02	rough	117.5	0 (very light)	Y	-0.14
16.05	smooth	112.2	4 (dark)	Y	-0.10
18.9	smooth	102.3	1 (light)	N	+0.21

# Prediction problems

- **Classification**

- E.g., predicting whether a treatment is successful vs. unsuccessful
- $Y$  is a finite discrete set (e.g., successful vs. unsuccessful treatment)

- **Regression**

- E.g., predicting the future size of a tumor
- $Y = \mathbb{R}$  (i.e., we are predicting a real number)

tumor size	texture	perimeter	shade	outcome	size change
18.02	rough	117.5	0 (very light)	Y	-0.14
16.05	smooth	112.2	4 (dark)	Y	-0.10
18.9	smooth	102.3	1 (light)	N	+0.21

# Variable types

- **Numerical**, real number measurements (usually *quantitative*).
  - Assumes that similar measurements are similar in nature.
- **Categorical**, from a discrete set (often *qualitative*)
  - E.g. {Spam, Not-spam}
- **Ordinal**, from a discrete set, without metric relation, but allows ranking.
  - E.g. {first, second, third}

tumor size	texture	perimeter	shade	outcome	size change
18.02	rough	117.5	0 (very light)	Y	-0.14
16.05	smooth	112.2	4 (dark)	Y	-0.10
18.9	smooth	102.3	1 (light)	N	+0.21

# The i.i.d. assumption

---

- In supervised learning, the examples  $\mathbf{x}_i$  in the training set are assumed to be **independently** and **identically distributed**.



# The i.i.d. assumption

---

- In supervised learning, the examples  $\mathbf{x}_i$  in the training set are assumed to be **independently** and **identically distributed**.
  - **Independently**: Every  $\mathbf{x}_i$  is freshly sampled according to some probability distribution  $\mathbf{D}$  over the data domain  $\mathbf{X}$ .
  - **Identically**: The distribution  $\mathbf{D}$  is the same for all examples.

Why?

# Empirical risk minimization

---

For a given function class  $\mathcal{F}$  and training sample  $S$ ,

- Define a notion of error (*left intentionally vague for now*):  
$$L_S(f) = \# \text{ mistakes made on the sample } S$$

# Empirical risk minimization

---

For a given function class  $\mathbf{F}$  and training sample  $\mathbf{S}$ ,

- Define a notion of error (*left intentionally vague for now*):  
$$L_{\mathbf{S}}(\mathbf{f}) = \# \text{ mistakes made on the sample } \mathbf{S}$$

- Define the Empirical Risk Minimization (ERM):

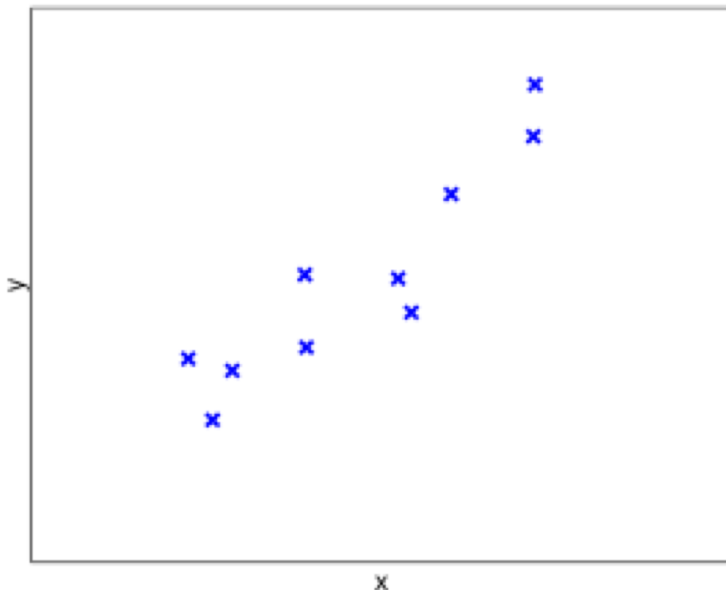
$$\text{ERM}_{\mathbf{F}}(\mathbf{S}) = \operatorname{argmin}_{\mathbf{f} \in \mathbf{F}} L_{\mathbf{S}}(\mathbf{f})$$

where **argmin** returns the function  $\mathbf{f}$  (or set of functions) that achieves the minimum loss on the training sample.

- Easier to minimize the error with i.i.d. assumption

# Empirical risk minimization

- What hypothesis class should we pick?



<i>Observe</i>	<i>Predict</i>
<u>x</u>	<u>y</u>
0.86	2.49
0.09	0.83
-0.85	-0.25
0.87	3.10
-0.44	0.87
-0.43	0.02
-1.1	-0.12
0.40	1.81
-0.96	-0.83
0.17	0.43

# Linear hypothesis

---

- Suppose  $y$  is a linear function of  $\mathbf{x}$ :

$$\begin{aligned}f_{\mathbf{w}}(\mathbf{x}) &= w_0 + w_1 x_1 + \dots + w_m x_m \\&= w_0 + \sum_{j=1:m} w_j x_j\end{aligned}$$

- The  $w_j$  are called parameters or weights.
- To simplify notation, we add an attribute  $x_0=1$  to the  $m$  other attributes (also called bias term or intercept).

# Linear hypothesis

---

- Suppose  $y$  is a linear function of  $\mathbf{x}$ :

$$\begin{aligned}f_{\mathbf{w}}(\mathbf{x}) &= w_0 + w_1 x_1 + \dots + w_m x_m \\&= w_0 + \sum_{j=1:m} w_j x_j\end{aligned}$$

- The  $w_j$  are called parameters or weights.
- To simplify notation, we add an attribute  $x_0=1$  to the  $m$  other attributes (also called bias term or intercept).

**How should we pick the *weights*?**

# Least-squares solution method

---

- The linear regression problem:  $f_{\mathbf{w}}(\mathbf{x}) = w_0 + \sum_{j=1:m} w_j x_j$   
where  $m$  = the dimension of observation space, i.e. number of features.
- **Goal:** Find the **best** linear model (i.e., weights) given the data.
- Many different possible **evaluation** criteria!
- Most common choice is to find the  $\mathbf{w}$  that minimizes:

$$\text{Err}(\mathbf{w}) = \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

(A note on notation: Technically  $\mathbf{w}$  and  $\mathbf{x}$  are vectors of size  $m+1$ , i.e., number of features + the dummy/intercept term. However, for notational simplicity from now on we say that  $\mathbf{w}$  and  $\mathbf{x}$  are vectors of size  $m$ , with  $m$  = “number of features + 1” and just treat the intercept as an extra feature)

# Least-squares solution method

- Goal: Find a function of the form  $f_w(\mathbf{x}) = w_0 + \sum_{j=1:m} w_j x_j$
- such that

$$f_w = \operatorname{argmin} \sum_{i=1:n} (y_i - w^T x_i)^2$$

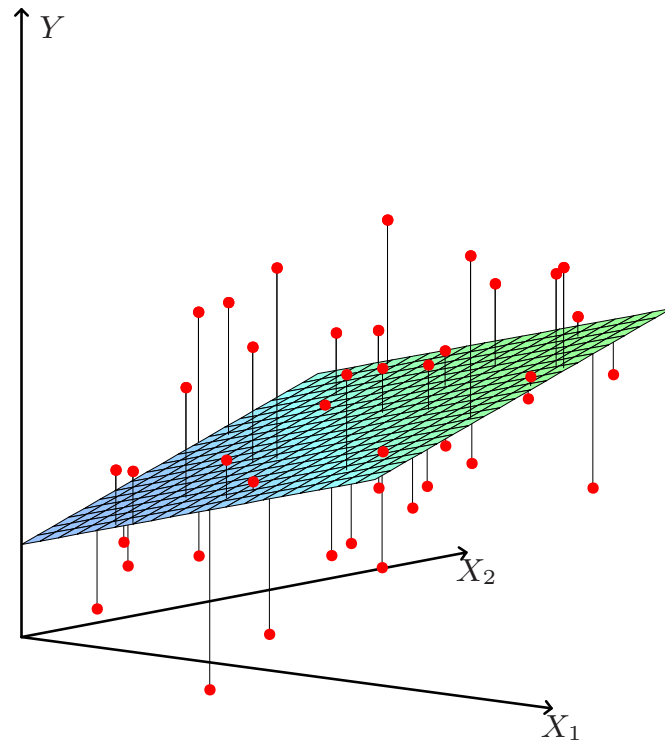
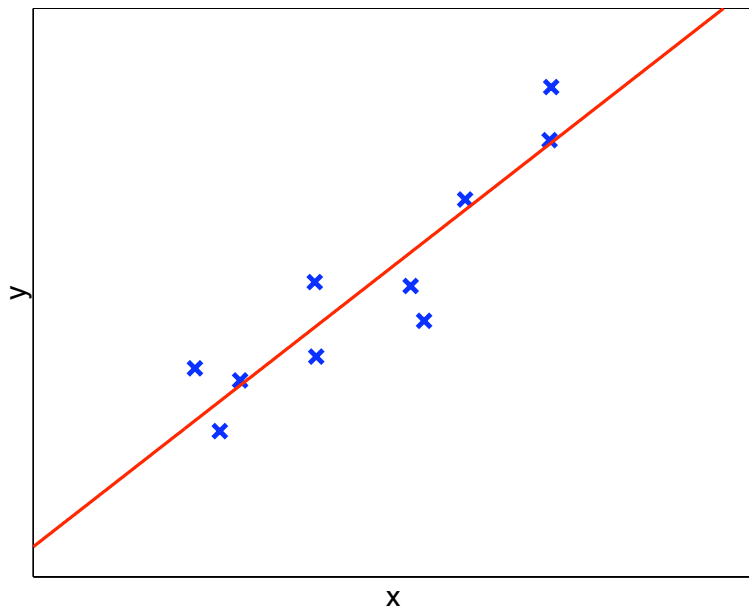
This is “least-squares”  
solution

True target value

Predicted value for  
the target



# Least-squares solution



# Least-squares solution method

---

- Re-write in matrix notation:  $f_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w}$   
$$\text{Err}(\mathbf{w}) = (\mathbf{y} - \mathbf{X}\mathbf{w})^T(\mathbf{y} - \mathbf{X}\mathbf{w})$$

where

$\mathbf{X}$  is the  $n \times m$  matrix of input data,  
 $\mathbf{y}$  is the  $n \times 1$  vector of output data,  
 $\mathbf{w}$  is the  $m \times 1$  vector of weights.

# Least-squares solution method

---

- Re-write in matrix notation:  $f_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w}$   
$$\text{Err}(\mathbf{w}) = (\mathbf{y} - \mathbf{X}\mathbf{w})^T(\mathbf{y} - \mathbf{X}\mathbf{w})$$

where  $\mathbf{X}$  is the  $n \times m$  matrix of input data,  
 $\mathbf{y}$  is the  $n \times 1$  vector of output data,  
 $\mathbf{w}$  is the  $m \times 1$  vector of weights.

- To minimize, take the derivative w.r.t.  $\mathbf{w}$ :
$$\partial \text{Err}(\mathbf{w}) / \partial \mathbf{w} = -2 \mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$
  - You get a system of  $m$  equations with  $m$  unknowns.

# Least-squares solution method

- Re-write in matrix notation:  $f_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w}$   
$$\text{Err}(\mathbf{w}) = (\mathbf{y} - \mathbf{X}\mathbf{w})^T(\mathbf{y} - \mathbf{X}\mathbf{w})$$

where  $\mathbf{X}$  is the  $n \times m$  matrix of input data,  
 $\mathbf{y}$  is the  $n \times 1$  vector of output data,  
 $\mathbf{w}$  is the  $m \times 1$  vector of weights.

- To minimize, take the derivative w.r.t.  $\mathbf{w}$ :  
$$\partial \text{Err}(\mathbf{w}) / \partial \mathbf{w} = -2 \mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$
  - You get a system of  $m$  equations with  $m$  unknowns.
- Set these equations to 0: 
$$\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}) = 0$$

# Least-squares solution method

---

- We want to solve for  $\mathbf{w}$ :  $\mathbf{X}^T (\mathbf{Y} - \mathbf{X}\mathbf{w}) = 0$
- Try a little algebra:  
 $\mathbf{X}^T \mathbf{Y} = \mathbf{X}^T \mathbf{X} \mathbf{w}$   
 $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$   
( $\hat{\mathbf{w}}$  denotes the estimated weights)

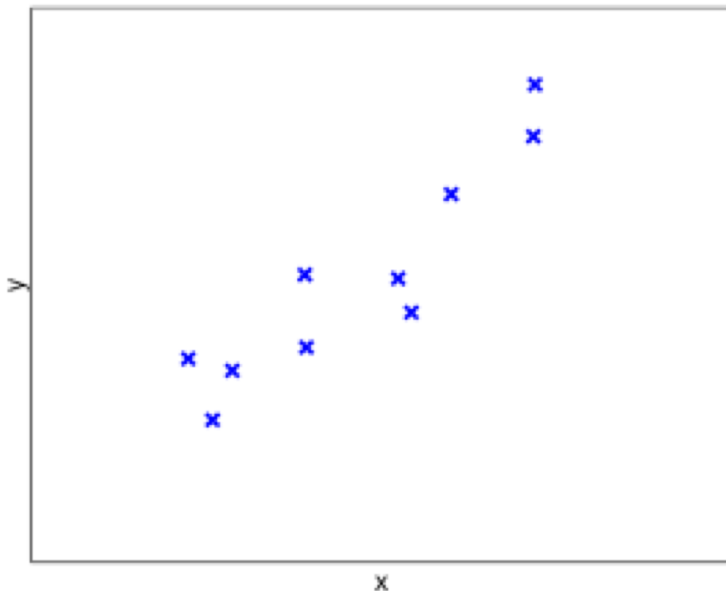
# Least-squares solution method

---

- We want to solve for  $\mathbf{w}$ :  $\mathbf{X}^T (\mathbf{Y} - \mathbf{X}\mathbf{w}) = 0$
- Try a little algebra:  
 $\mathbf{X}^T \mathbf{Y} = \mathbf{X}^T \mathbf{X} \mathbf{w}$   
 $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$   
( $\hat{\mathbf{w}}$  denotes the estimated weights)
- The fitted data:  $\hat{\mathbf{Y}} = \mathbf{X}\hat{\mathbf{w}} = \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$
- To predict new data  $\mathbf{X}' \rightarrow \mathbf{Y}'$ :  $\mathbf{Y}' = \mathbf{X}'\hat{\mathbf{w}} = \mathbf{X}' (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$

# Example of linear regression

- What is a plausible estimate of  $w$ ?



<i>Observe</i>	<i>Predict</i>
<u>x</u>	<u>y</u>
0.86	2.49
0.09	0.83
-0.85	-0.25
0.87	3.10
-0.44	0.87
-0.43	0.02
-1.1	-0.12
0.40	1.81
-0.96	-0.83
0.17	0.43

# Example of linear regression

- What is a plausible estimate of  $w$ ?
- Recall that our **least-squares estimate** is

- $$\hat{w} = (X^T X)^{-1} X^T Y$$

<i>Observe</i>	<i>Predict</i>
<u>x</u>	<u>y</u>
0.86	2.49
0.09	0.83
-0.85	-0.25
0.87	3.10
-0.44	0.87
-0.43	0.02
-1.1	-0.12
0.40	1.81
-0.96	-0.83
0.17	0.43



# Data matrices

---

$$\begin{aligned} X^T X &= \begin{bmatrix} 0.86 & 0.09 & -0.85 & 0.87 & -0.44 & -0.43 & -1.10 & 0.40 & -0.96 & 0.17 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix} \times \begin{bmatrix} 0.86 & 1 \\ 0.09 & 1 \\ -0.85 & 1 \\ 0.87 & 1 \\ -0.44 & 1 \\ -0.43 & 1 \\ -1.10 & 1 \\ 0.40 & 1 \\ -0.96 & 1 \\ 0.17 & 1 \end{bmatrix} \\ &= \begin{bmatrix} 4.95 & -1.39 \\ -1.39 & 10 \end{bmatrix} \end{aligned}$$

# Data matrices

---

$$\begin{aligned} X^T Y &= \begin{bmatrix} 0.86 & 0.09 & -0.85 & 0.87 & -0.44 & -0.43 & -1.10 & 0.40 & -0.96 & 0.17 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix} \times \begin{bmatrix} 2.49 \\ 0.83 \\ -0.25 \\ 3.10 \\ 0.87 \\ 0.02 \\ -0.12 \\ 1.81 \\ -0.83 \\ 0.43 \end{bmatrix} \\ &= \begin{bmatrix} 6.49 \\ 8.34 \end{bmatrix} \end{aligned}$$

# Solving the problem

---

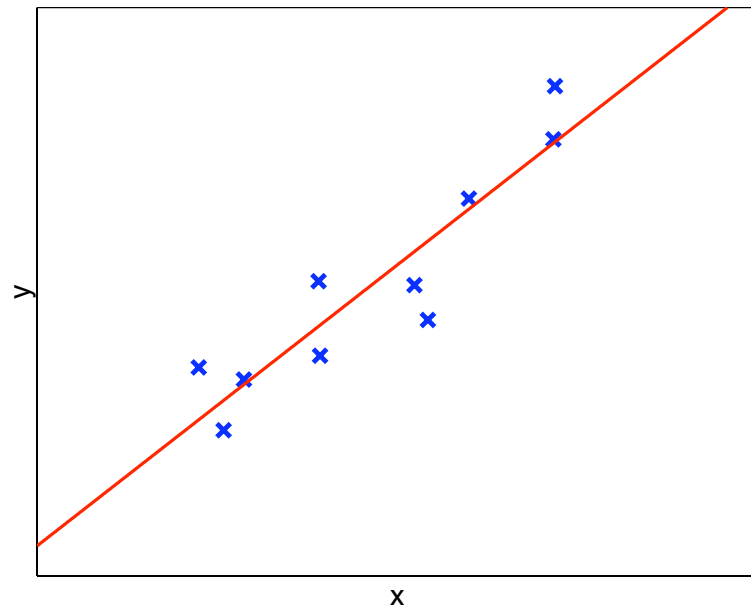
$$\mathbf{w} = (X^T X)^{-1} X^T Y = \begin{bmatrix} 4.95 & -1.39 \\ -1.39 & 10 \end{bmatrix}^{-1} \begin{bmatrix} 6.49 \\ 8.34 \end{bmatrix} = \begin{bmatrix} 1.60 \\ 1.05 \end{bmatrix}$$

So the best fit line is  $y = 1.60x + 1.05$ .

# Solving the problem

$$\mathbf{w} = (X^T X)^{-1} X^T Y = \begin{bmatrix} 4.95 & -1.39 \\ -1.39 & 10 \end{bmatrix}^{-1} \begin{bmatrix} 6.49 \\ 8.34 \end{bmatrix} = \begin{bmatrix} 1.60 \\ 1.05 \end{bmatrix}$$

So the best fit line is  $y = 1.60x + 1.05$ .



# Least-squares solution method

- Linear fit for a prostate cancer dataset
  - Features  $X = \{\text{lcavol}, \text{lweight}, \text{age}, \text{lbph}, \text{svi}, \text{lcp}, \text{gleason}, \text{pgg45}\}$
  - Output  $y$  = level of PSA (an enzyme which is elevated with cancer)..

Coefficient (i.e., learned weight,  $w_i$ )

- How does increasing  $x_i$  change the output  $y_i$ ?

Standard error

- How confident/precise is the estimate of  $w_i$ ?
- How “predictive” of  $y_i$  is  $x_i$ ?

Term	Coefficient	Std. Error
Intercept	$w_0 = 2.46$	0.09
lcavol	0.68	0.13
lweight	0.26	0.10
age	-0.14	0.10
lbph	0.21	0.10
svi	0.31	0.12
lcp	-0.29	0.15
gleason	-0.02	0.15
pgg45	0.27	0.15

# Least-squares solution method

- Linear fit for a prostate cancer dataset
  - Features  $X = \{\text{lcavol}, \text{lweight}, \text{age}, \text{lbph}, \text{svi}, \text{lcp}, \text{gleason}, \text{pgg45}\}$
  - Output  $y$  = level of PSA (an enzyme which is elevated with cancer)..

Large **coefficient** and  
small **standard error**



Expect a small change in  $x_j$   
to have a large effect on  $y$

Term	Coefficient	Std. Error
Intercept	$w_0 = 2.46$	0.09
lcavol	0.68	0.13
lweight	0.26	0.10
age	-0.14	0.10
lbph	0.21	0.10
svi	0.31	0.12
lcp	-0.29	0.15
gleason	-0.02	0.15
pgg45	0.27	0.15

# Computational cost of linear regression

---

- What operations are necessary?

# Computational cost of linear regression

---

- What operations are necessary?
  - Overall: 1 matrix inversion + 3 matrix multiplications
  - $\mathbf{X}^T \mathbf{X}$  (other matrix multiplications require fewer operations.)
    - $\mathbf{X}^T$  is  $m \times n$  and  $\mathbf{X}$  is  $n \times m$ , so we need  $nm^2$  operations.
  - $(\mathbf{X}^T \mathbf{X})^{-1}$ 
    - $\mathbf{X}^T \mathbf{X}$  is  $m \times m$ , so we need  $m^3$  operations.



# Computational cost of linear regression

---

- What operations are necessary?
  - Overall: 1 matrix inversion + 3 matrix multiplications
  - $\mathbf{X}^T\mathbf{X}$  (other matrix multiplications require fewer operations.)
    - $\mathbf{X}^T$  is  $m \times n$  and  $\mathbf{X}$  is  $n \times m$ , so we need  $nm^2$  operations.
  - $(\mathbf{X}^T\mathbf{X})^{-1}$ 
    - $\mathbf{X}^T\mathbf{X}$  is  $m \times m$ , so we need  $m^3$  operations.
- Overall, we have  $O(m^3 + nm^2)$  (i.e, polynomial) computational cost, but handling really large datasets (e.g, many points or features) can still be an issue!

# A more efficient alternative?

- Recall the least-squares solution:  $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$
- What if  $\mathbf{X}$  is too big to compute this explicitly (e.g.  $m \sim 10^6$ )?
- Go back to the gradient step:  $\text{Err}(\mathbf{w}) = (\mathbf{Y} - \mathbf{X}\mathbf{w})^T (\mathbf{Y} - \mathbf{X}\mathbf{w})$

Gradient tells us how to move the parameters to maximally increase/decrease the error!



$$\begin{aligned}\partial \text{Err}(\mathbf{w}) / \partial \mathbf{w} &= -2 \mathbf{X}^T (\mathbf{Y} - \mathbf{X}\mathbf{w}) \\ &= 2(\mathbf{X}^T \mathbf{X}\mathbf{w} - \mathbf{X}^T \mathbf{Y})\end{aligned}$$

# Gradient descent for linear regression

---

- We want to produce a sequence of weight solutions,  $\mathbf{w}_0, \mathbf{w}_1, \mathbf{w}_2, \dots$ , such that:  $\text{Err}(\mathbf{w}_0) > \text{Err}(\mathbf{w}_1) > \text{Err}(\mathbf{w}_2) > \dots$

# Gradient descent for linear regression

- We want to produce a sequence of weight solutions,  $\mathbf{w}_0, \mathbf{w}_1, \mathbf{w}_2, \dots$ , such that:  $\text{Err}(\mathbf{w}_0) > \text{Err}(\mathbf{w}_1) > \text{Err}(\mathbf{w}_2) > \dots$
- The algorithm:     *Given an initial weight vector  $\mathbf{w}_0$ ,*  
                          *Do for  $k=1, 2, \dots$*   

$$\mathbf{w}_{k+1} = \mathbf{w}_k - \alpha_k \partial \text{Err}(\mathbf{w}_k) / \partial \mathbf{w}_k$$

  
                          *End when  $|\mathbf{w}_{k+1} - \mathbf{w}_k| < \varepsilon$*

Take a “step” in the (negative) direction specified by the gradient.

# Gradient descent for linear regression

---

- We want to produce a sequence of weight solutions,  $\mathbf{w}_0, \mathbf{w}_1, \mathbf{w}_2, \dots$ , such that:  $\text{Err}(\mathbf{w}_0) > \text{Err}(\mathbf{w}_1) > \text{Err}(\mathbf{w}_2) > \dots$
- The algorithm:     *Given an initial weight vector  $\mathbf{w}_0$ ,*  
                          *Do for  $k=1, 2, \dots$*   
                                  
$$\mathbf{w}_{k+1} = \mathbf{w}_k - \alpha_k \partial \text{Err}(\mathbf{w}_k) / \partial \mathbf{w}_k$$
  
                          *End when  $\|\mathbf{w}_{k+1} - \mathbf{w}_k\| < \varepsilon$*
- Parameter  $\alpha_k > 0$  is the step-size (or learning rate) for iteration  $k$ .

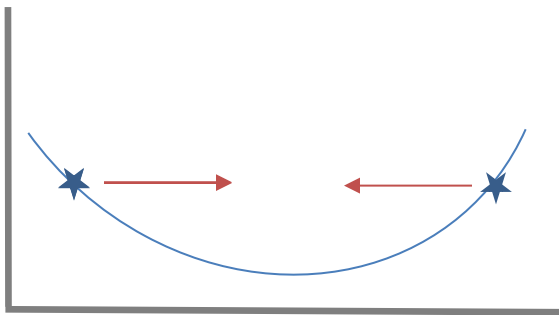
# Convergence

---

- Convergence depends in part on the  $\alpha_k$ .
- If steps are too large: the  $\mathbf{w}_k$  may oscillate forever.
  - This suggests that  $\alpha_k \rightarrow 0$  as  $k \rightarrow \infty$ .
- If steps are too small: the  $\mathbf{w}_k$  may not move far enough to reach a local minimum.

# Convergence

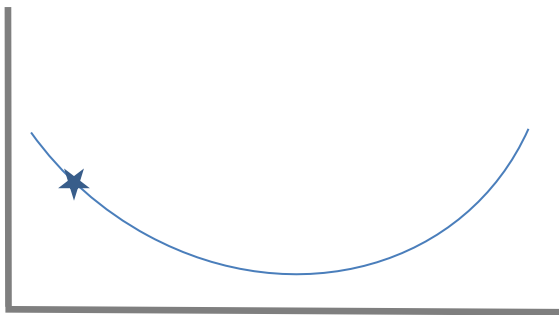
- Convergence depends in part on the  $\alpha_k$ .
- If steps are too large: the  $\mathbf{w}_k$  may oscillate forever.
  - This suggests that  $\alpha_k \rightarrow 0$  as  $k \rightarrow \infty$ .



# Convergence

---

- Convergence depends in part on the  $\alpha_k$ .
- If steps are too small: the  $\mathbf{w}_k$  may not move far enough to reach a local minimum.





# Robbins-Monroe conditions

---

- The  $\alpha_k$  are a Robbins-Monroe sequence if:

$$\sum_{k=0:\infty} \alpha_k = \infty$$

$$\sum_{k=0:\infty} \alpha_k^2 < \infty$$

- These conditions are sufficient to ensure convergence of the  $\mathbf{w}_k$  to a local minimum of the error function.

# Robbins-Monroe conditions

---

- The  $\alpha_k$  are a Robbins-Monroe sequence if:

$$\sum_{k=0:\infty} \alpha_k = \infty$$

$$\sum_{k=0:\infty} \alpha_k^2 < \infty$$

- These conditions are sufficient to ensure convergence of the  $\mathbf{w}_k$  to a local minimum of the error function.

E.g.  $\alpha_k = 1 / (k + 1)$

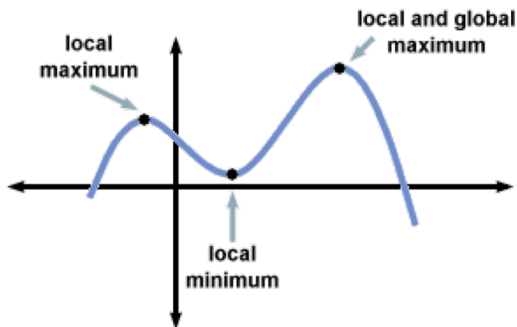
E.g.  $\alpha_k = 1/2$  for  $k = 1, \dots, T$

$$\alpha_k = 1/2^2 \text{ for } k = T+1, \dots, (T+1)+2T$$

etc.

# Local vs. global minima

- Convergence is NOT to a global minimum, only to local minimum.

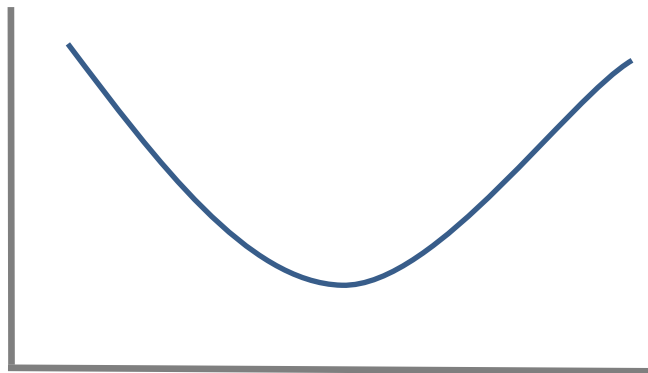


Only an issue for “non-convex” functions. For “convex” functions (e.g., simple linear regression) all local minima are also global minima.

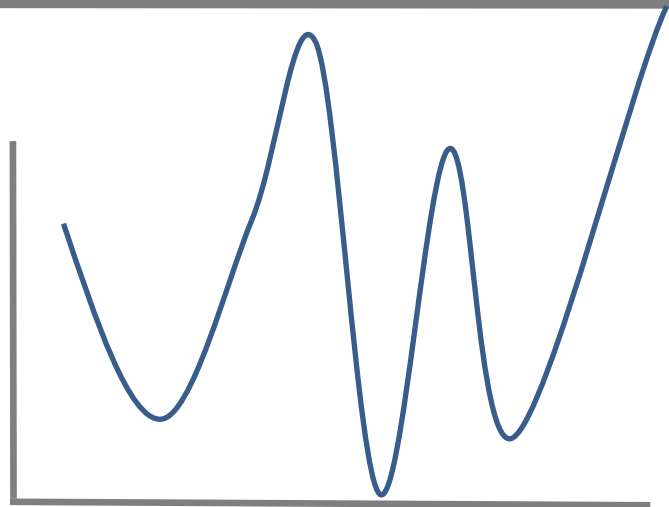
- The blue line represents the **error function**. There is no guarantee regarding the amount of error of the weight vector found by gradient descent, compared to the globally optimal solution. (Random restarts can help.)

# Convex vs. non-convex (informally)

---



Convex

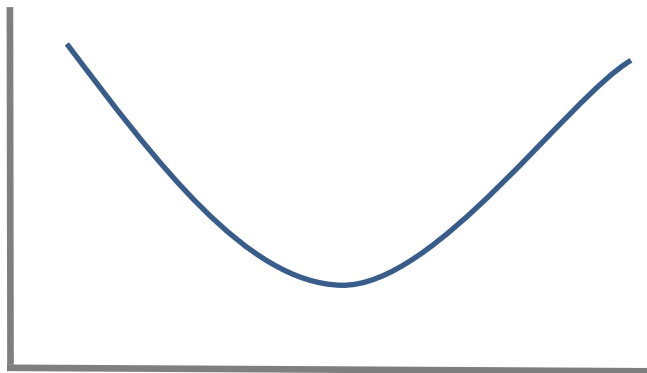


Non-convex

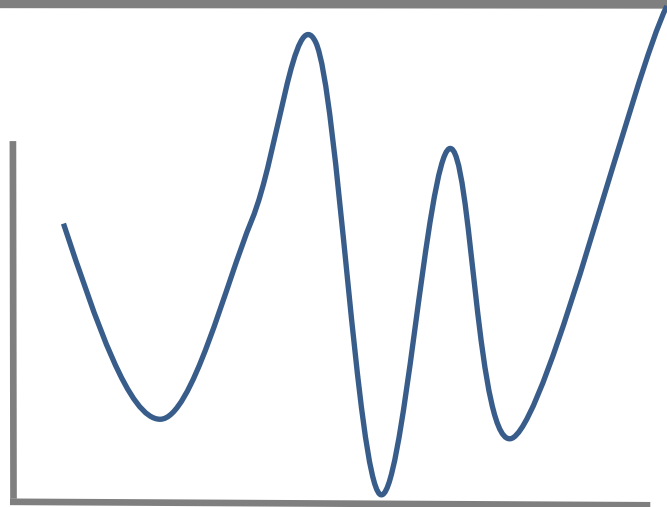
Most state-of-the-art approaches (e.g., deep learning) are non-convex!

# Convex vs. non-convex (informally)

---



Convex



Non-convex

Most state-of-the-art approaches (e.g., deep learning) are non-convex!

But linear regression is convex!

# What you should know

---

- Definition and characteristics of a supervised learning problem.
- Linear regression (hypothesis class, error function, algorithm).
- Closed-form least-squares solution method (algorithm, computational complexity, stability issues).
- Gradient descent method (algorithm, properties).

# To-do list

---

- Reproduce the linear regression example (slides 22-27), solving it using the software of your choice.
- Practice/self-assessment quizzes this week. You should be able to get 100% on these on the second try!
- Suggested complementary readings:
  - Ch.2 (Sec. 2.1-2.4, 2.9) of Hastie et al.
  - Ch.3 of Bishop.
  - Ch.9 of Shalev-Schwartz et al.
- Write down the midterm date/time (November 18th, 6-8pm) and contact the head TA (Joey Bose, [joey.bose@mail.mcgill.ca](mailto:joey.bose@mail.mcgill.ca)) right away if you know you need to be away at this time.